

neutron-lan

SDN study environnement @ home

github.com/alexanderplatz1999

Last update: May 27th, 2014

Background and Motivation

- My belief
 - SDN = Python and Java defines network
- Too many SDN definitions
 - I have been confused a lot.
 - OpenFlow, OVSD, Netconf, BGP extensions such as FlowSpec...
 - The latest addition: OpFlex (DevOps-like)
- What's the real SDN?
 - Let's develop SDN by myself and examine every definition.
- But, wait! I need a SDN study environment at home.
 - I am a poor guy, so I cannot buy expensive SDN-capable switches from Cisco, Juniper...

Strategy

- My budget is less than \$200.
- Switches/routers I purchased in Akihabara, Tokyo
 - Three \$40 broadband routers and one \$40 Raspberry Pi
- And I develop all the SDN software from scratch
 - But reuse existing networking software as much as possible, such as Open vSwitch
- Base knowledge/skills
 - SDN in the past: SIP and IP-PBX
 - OpenFlow, OpenStack neutron and SaltStack
 - Java and Python
 - HTML5 and CSS (a little)
- Let's develop neutron-like SDN for my home network ⇒ let's call it 'neutron-lan'

Project 'neutron-lan' characteristics


- Cheap routers as 'Baremetal Switch'
 - OpenWrt routers and Raspberry Pi
 - u-boot for installing new firmware
- Home-made DevOps tool 'NLAN' from scratch
 - 100% Python implementation
 - YAML-based state rendering
 - Model-driven service abstraction
- VXLAN-based edge-overlay for network virtualization
- LXC for Network Functions Virtualization
- Open vSwitch as a programmable switch
- OVSDB as a general-purpose config database

Project neutron-lan

Three major works so far (Dec/2013 ~ May/2014)

- **Rebuilding OpenWrt/Raspbian kernel/kernel-modules capable of OVS 2.0.0, network namespaces(netns), veth(virtual ether) and LXC.**
- **Configuring VXLAN-based edge-overlay for my home network**
- **Developing a home-made DevOps tool “NLAN”**

Cheap routers as 'Baremetal Switch'



Test bed (cont'd)

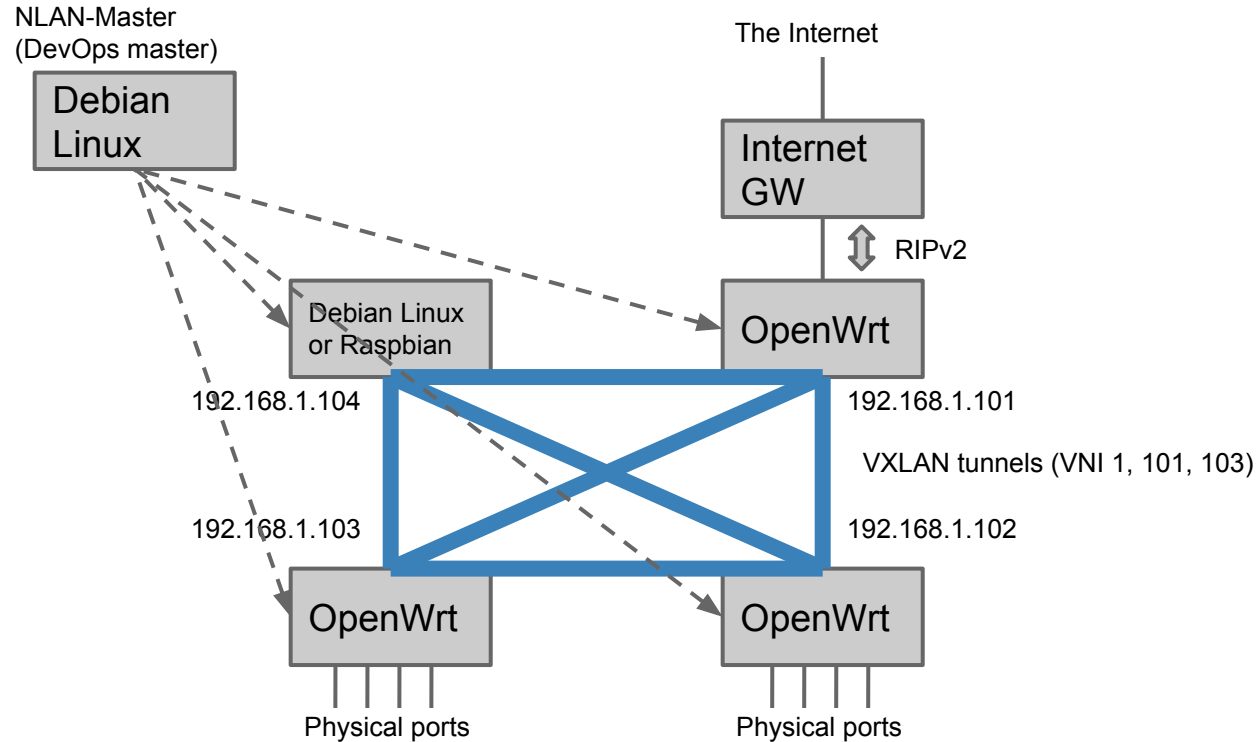


OpenWrt routers
(and Home Gateway)

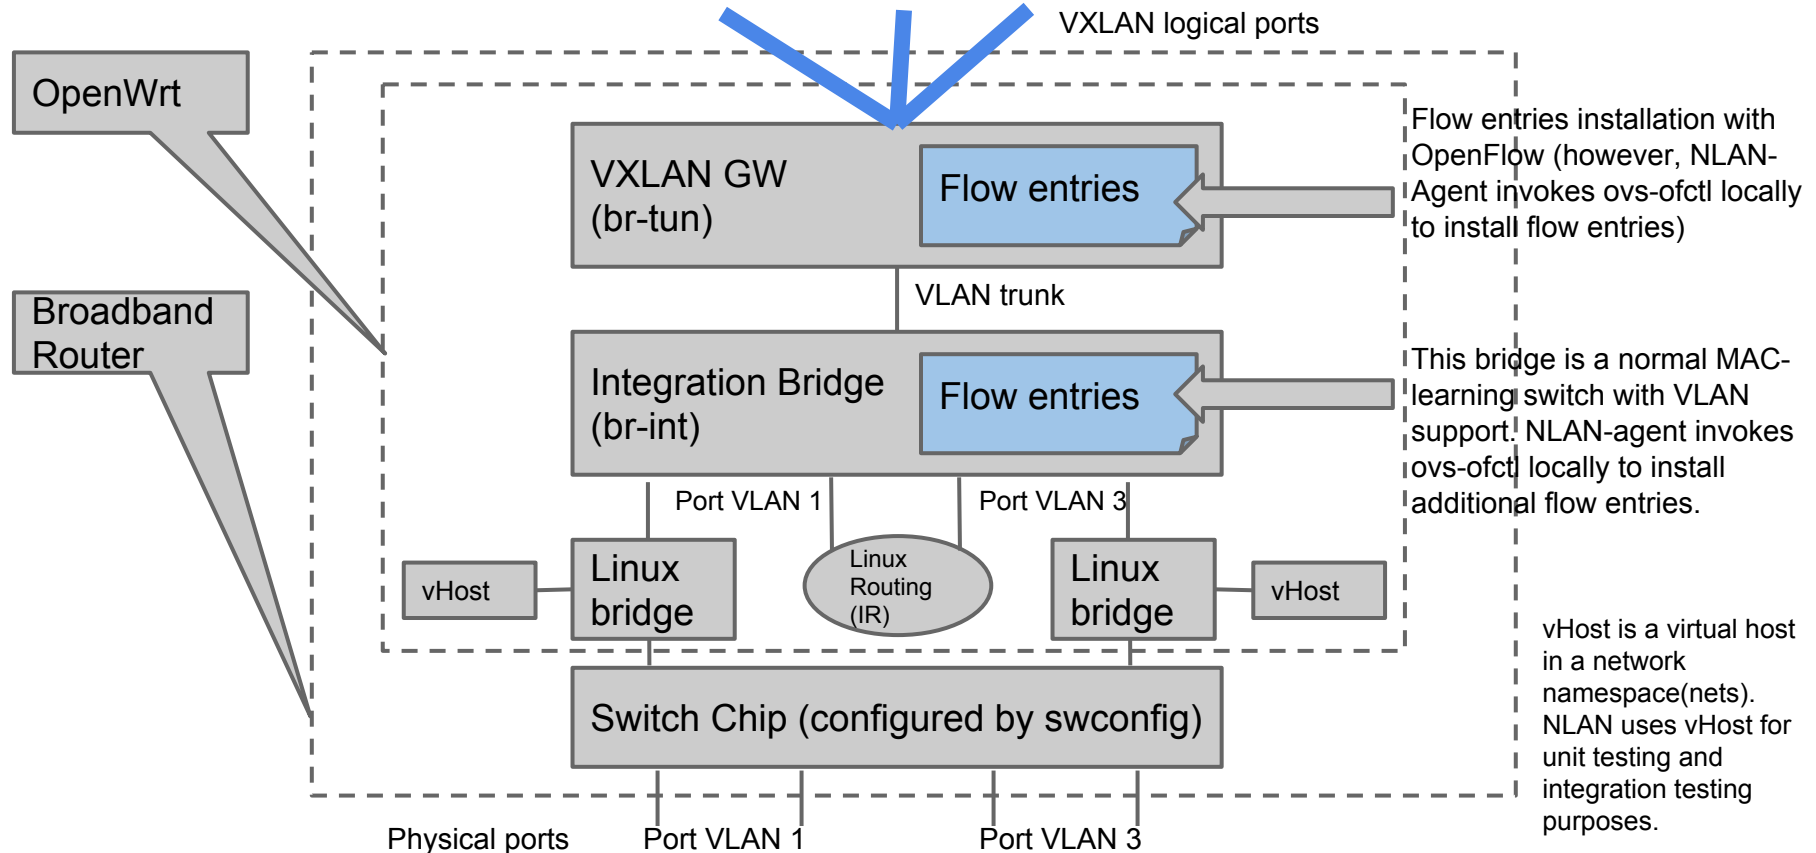


Raspberry Pi

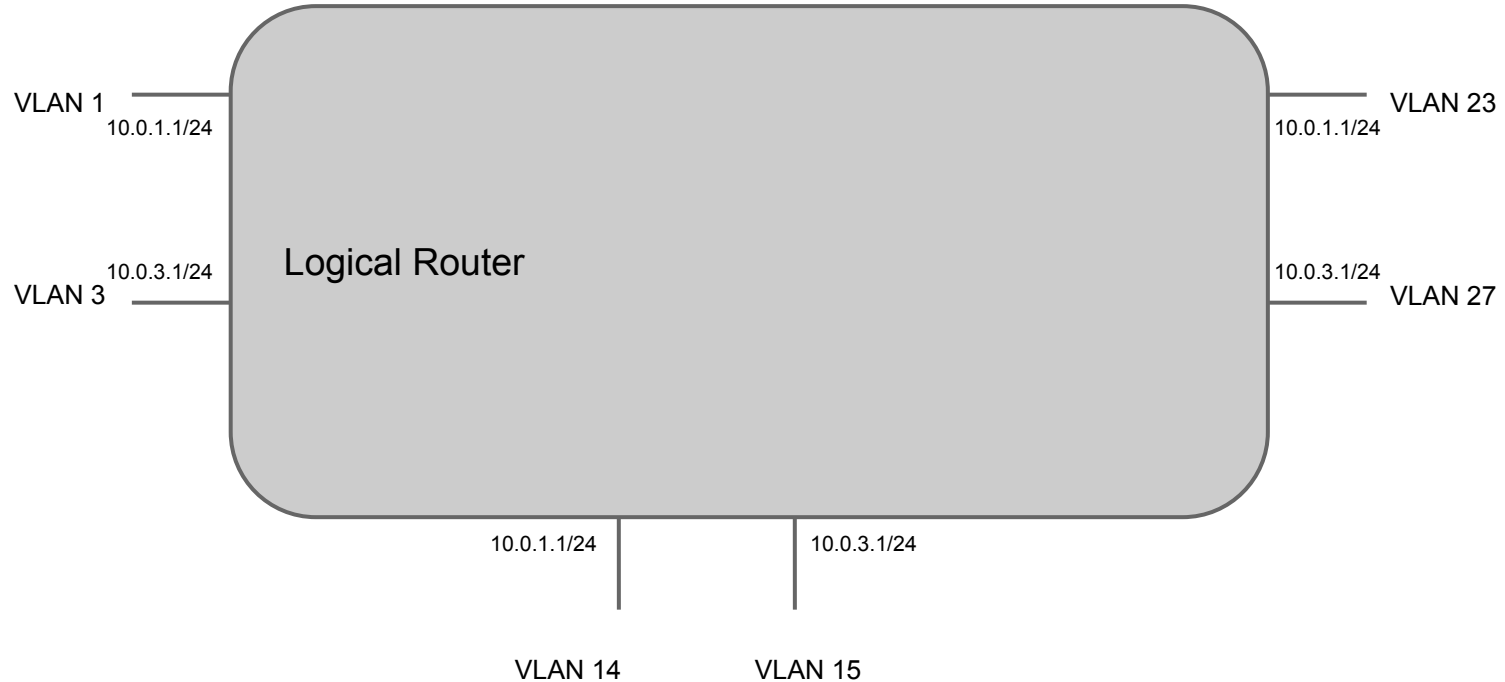
Test bed



OpenStack-neutron-like bridge configuration



Distributed Virtual Router (Logical view)



Virtual network topologies

NLAN node operation mode	Virtual Network Topology
dvr	Distributed Virtual Router
hub	Hub & Spoke
spoke	Hub & Spoke
spoke_dvr	Mixture of DVR and Hub & Spoke

NLAN “subnets” state and its parameters in YAML

subnets:

- vid: 1

vni: 1001

ip_dvr:

addr: '10.0.1.1/24'

mode: **hub**  mode can be 'dvr', 'hub', 'spoke' or 'spoke_dvr'

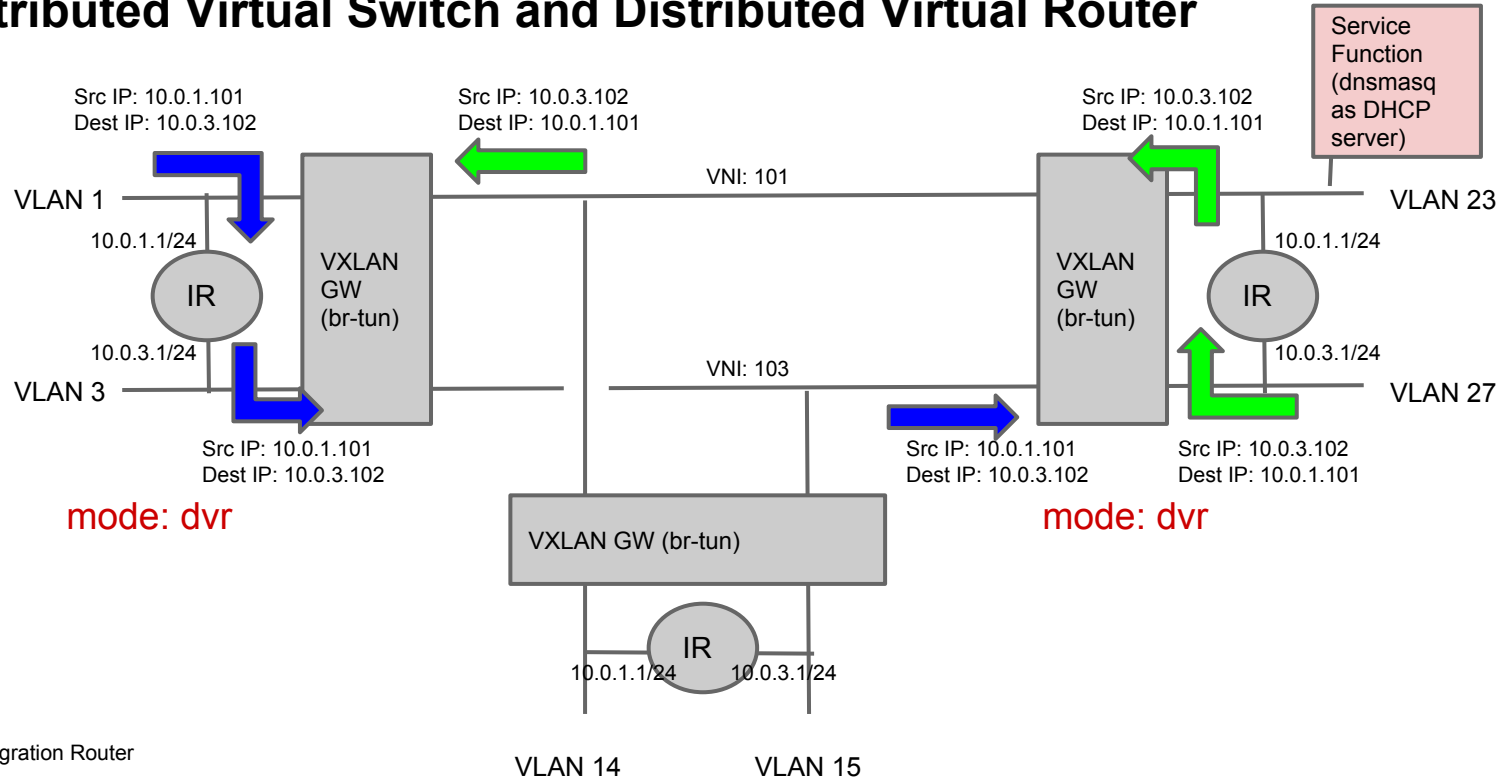
dhcp: **enabled**

ip_vhost: '10.0.1.101/24'

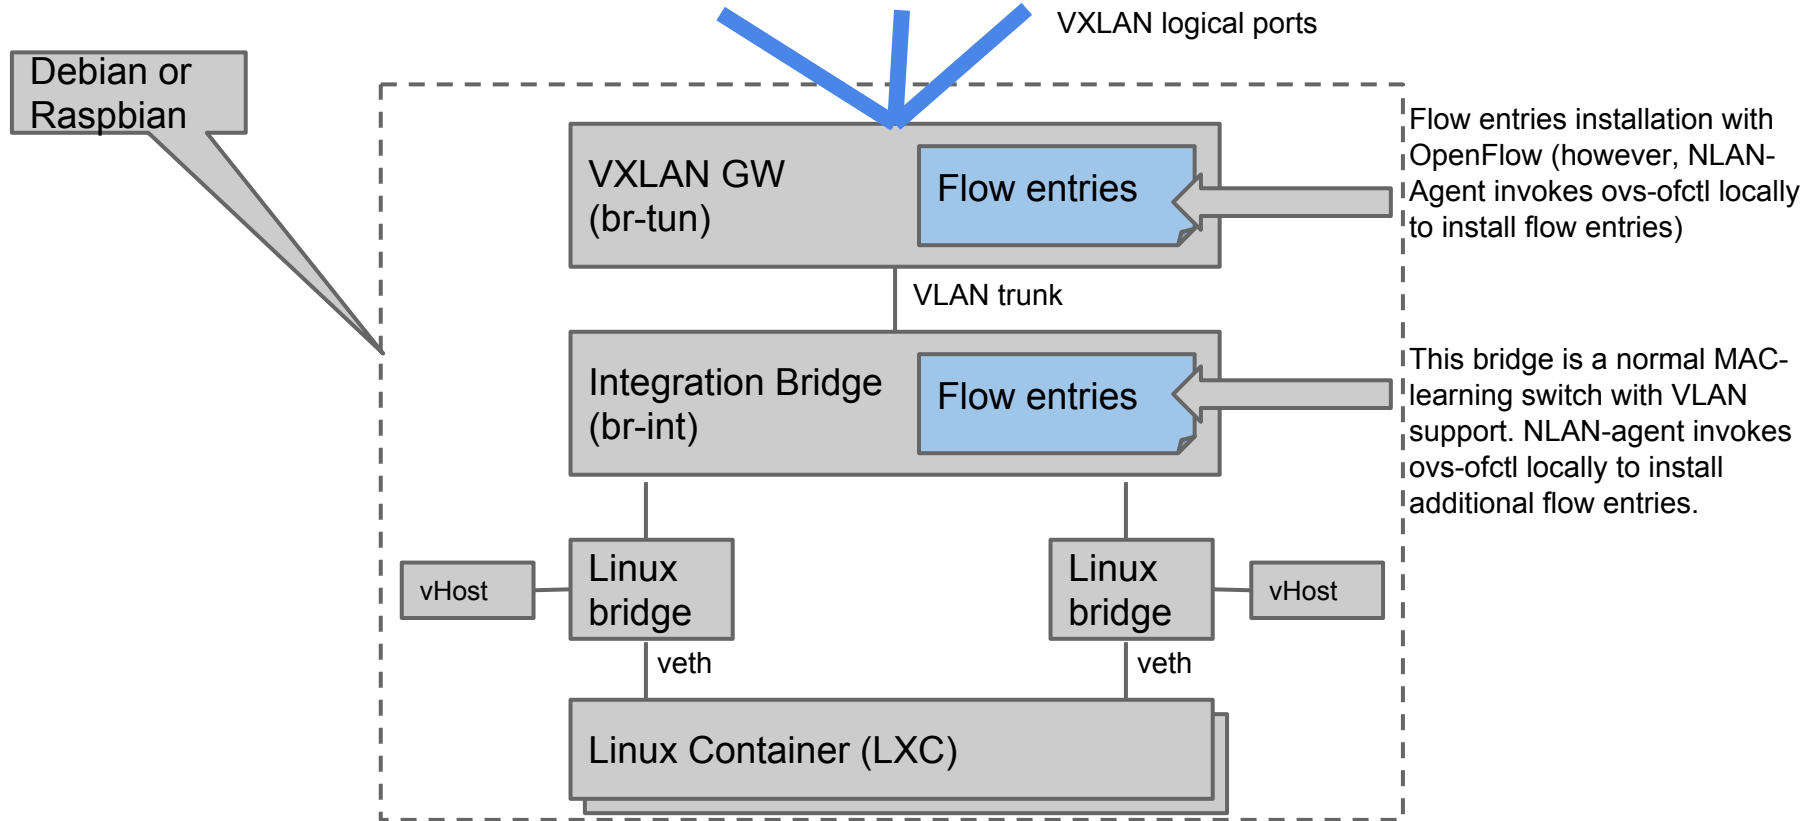
ports: [eth0.1]

peers: <peers>

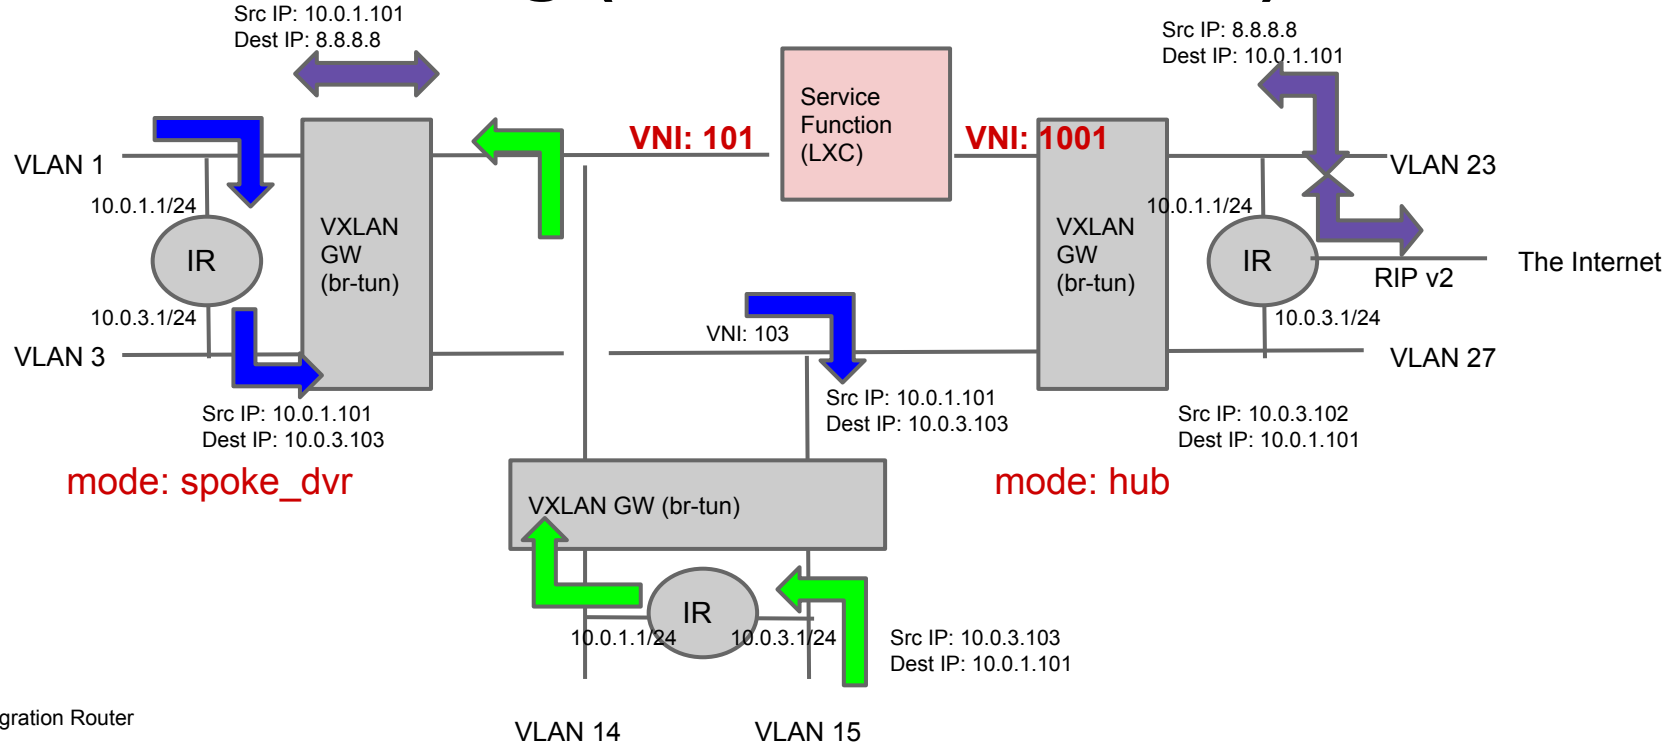
Distributed Virtual Switch and Distributed Virtual Router




Service Function in Linux Container



Service Chaining (Service Insertion)



Example: Snort (in IPS mode) as Service Function



```
04/13-23:16:32.859385 10.0.1.102 -> 8.8.8.8  
ICMP TTL:64 TOS:0x0 ID:11745 IpLen:20 DgmLen:84 DF  
Type:8 Code:0 ID:49496 Seq:25 ECHO  
==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+==+=+
```

```
04/13-23:16:32.861970 8.8.8.8 -> 10.0.1.102  
ICMP TTL:63 TOS:0x0 ID:38077 IpLen:20 DgmLen:84  
Type:0 Code:0 ID:49496 Seq:25 ECHO REPLY  
+++++
```

```
04/13-23:16:33.861151 10.0.1.102 -> 8.8.8.8
ICMP TTL:64 TOS:0x0 ID:11746 IpLen:20 DgmLen:84 DF
Type:8 Code:0 ID:49496 Seq:26 ECHO
=====
```

```
04/13-23:16:33.862906 8.8.8.8 -> 10.0.1.102
ICMP TTL:63 TOS:0x0 ID:38078 IpLen:20 DgmLen:84
Type:0 Code:0 ID:49496 Seq:26 ECHO REPLY
+++++
```

Snort is a free and open source IPS/IDS software.

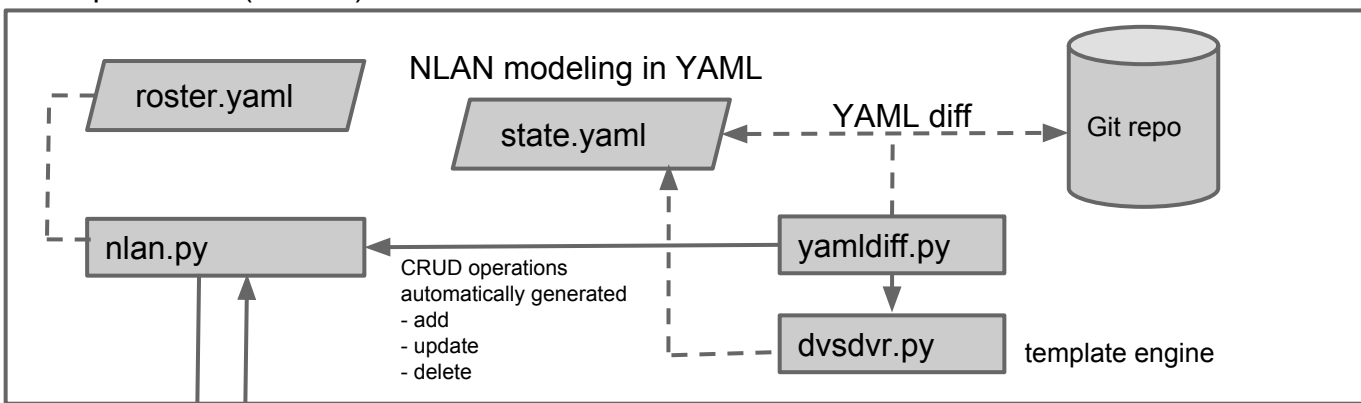
Home-made DevOps tool 'NLAN'

- 100% Python implementation
- Borrowed a lot of ideas from SaltStack
 - Model-driven approach
 - YAML-based state rendering w/ a simple template engine
 - Imperative/declarative state rendering
- Works with OpenWrt with minimal Python
 - opkg install python-mini
 - opkg install python-json
 - sshd
- OVSDB as a local config mgmt database
- State schema defined in YAML
 - merged with OVSDB schema in JSON

NLAN architecture (w/ config modules)

It's a bit like Salt State Modules...

DevOps master (Debian)

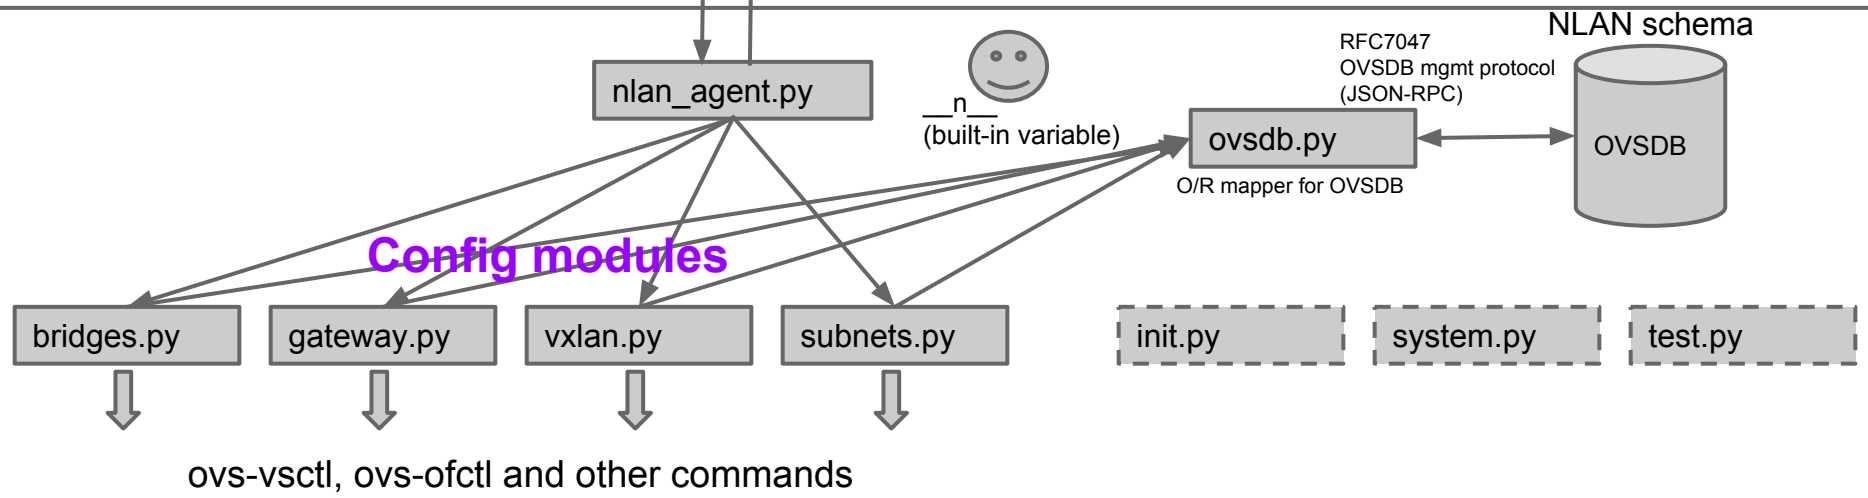


Python OrderedDict object
via STDIN

Command output
via STDOUT/STDERR

SSH

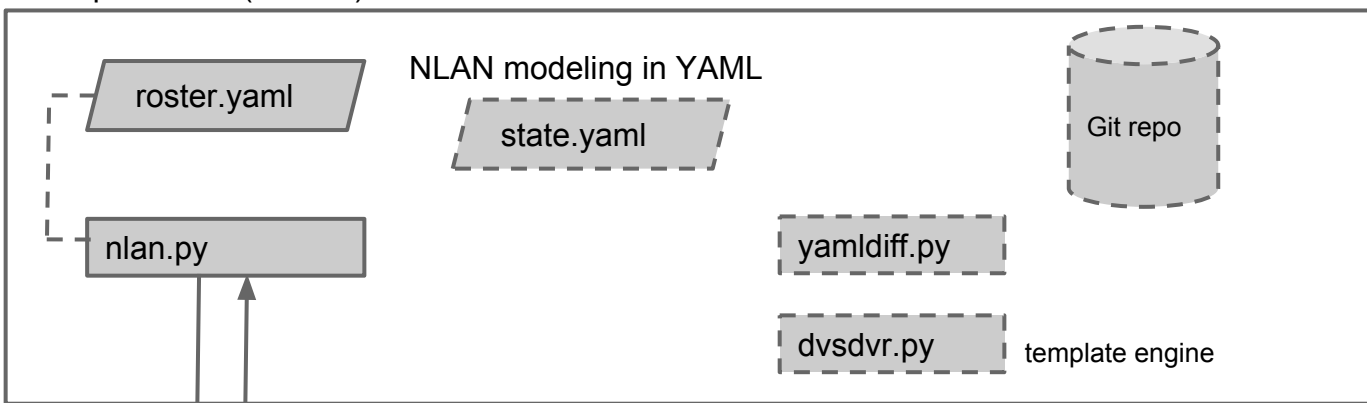
Router (OpenWrt or Raspbian)



NLAN architecture (w/ rpc modules)

It's a bit like Salt Execution Modules...

DevOps master (Debian)

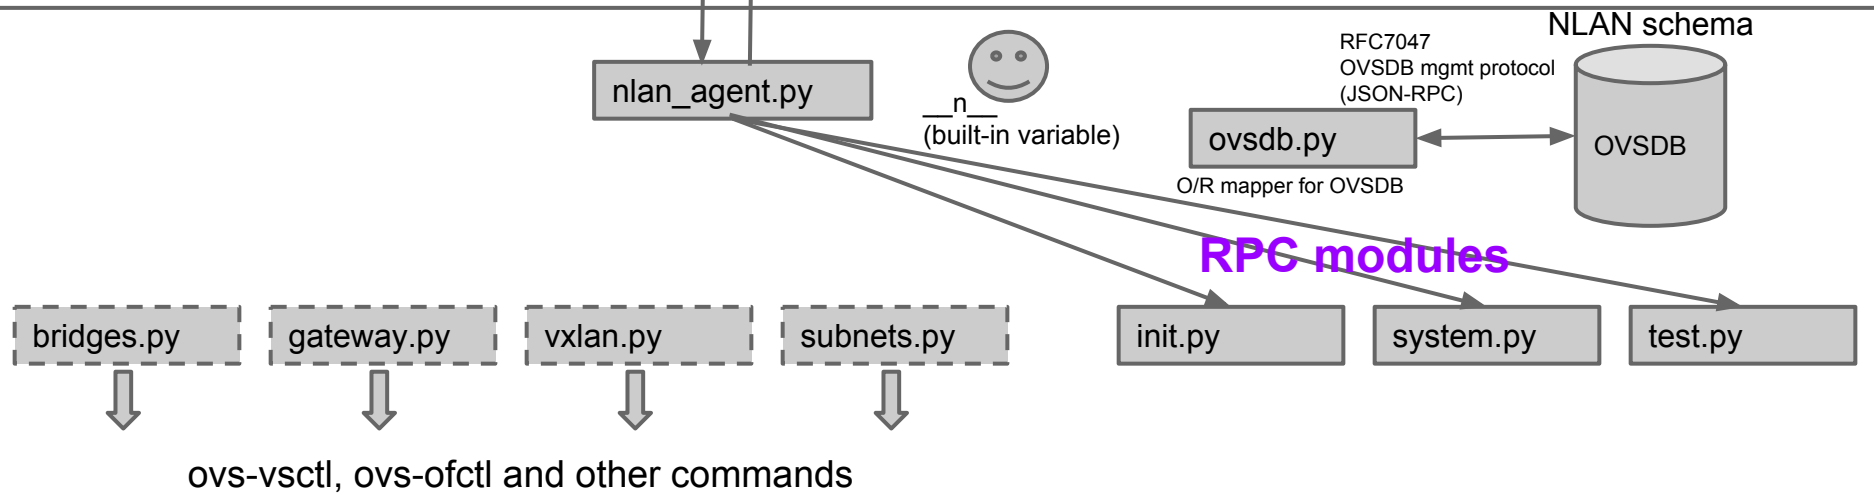


Python OrderedDict object
via STDIN

Command output
via STDOUT/STDERR

SSH

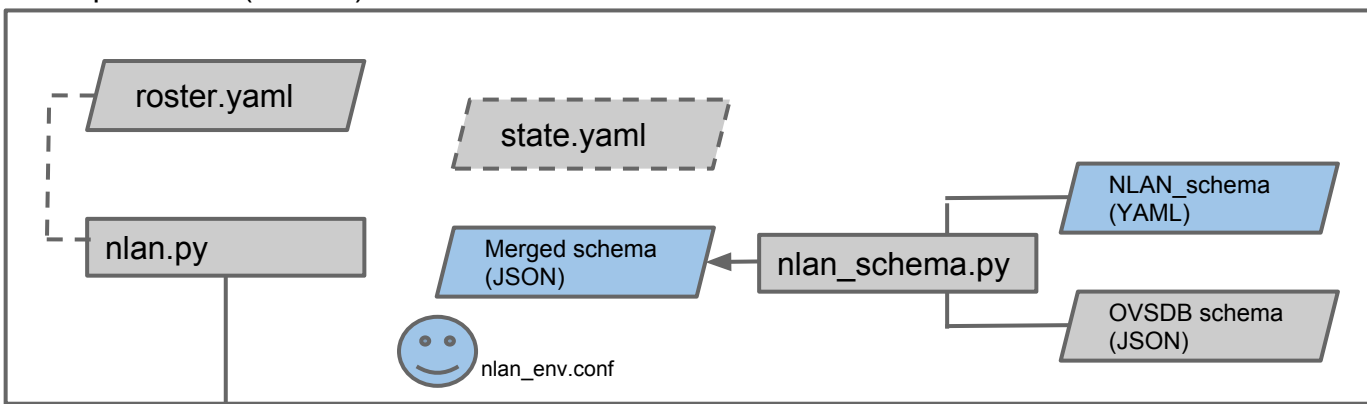
Router (OpenWrt or Raspbian)



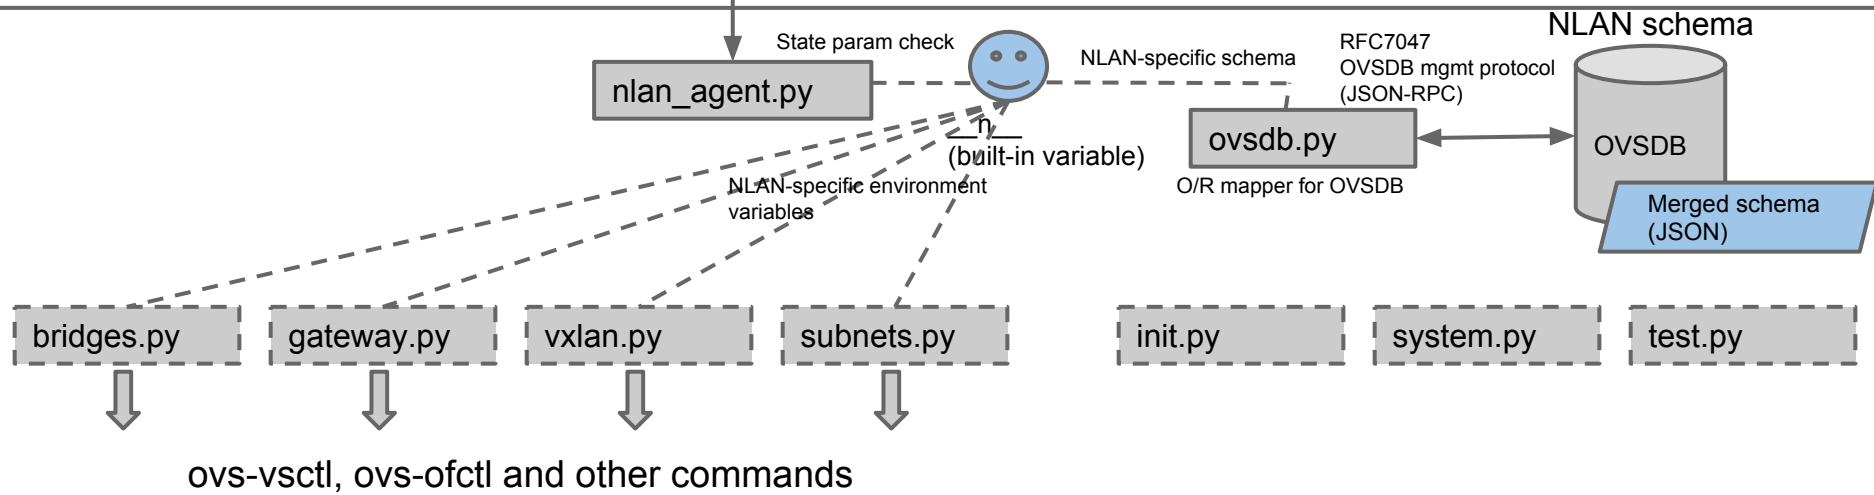
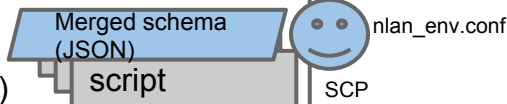
NLAN architecture (OVFlex)

Flexible OVSDb schemas (not so rigid)

DevOps master (Debian)



Router (OpenWrt or Raspbian)



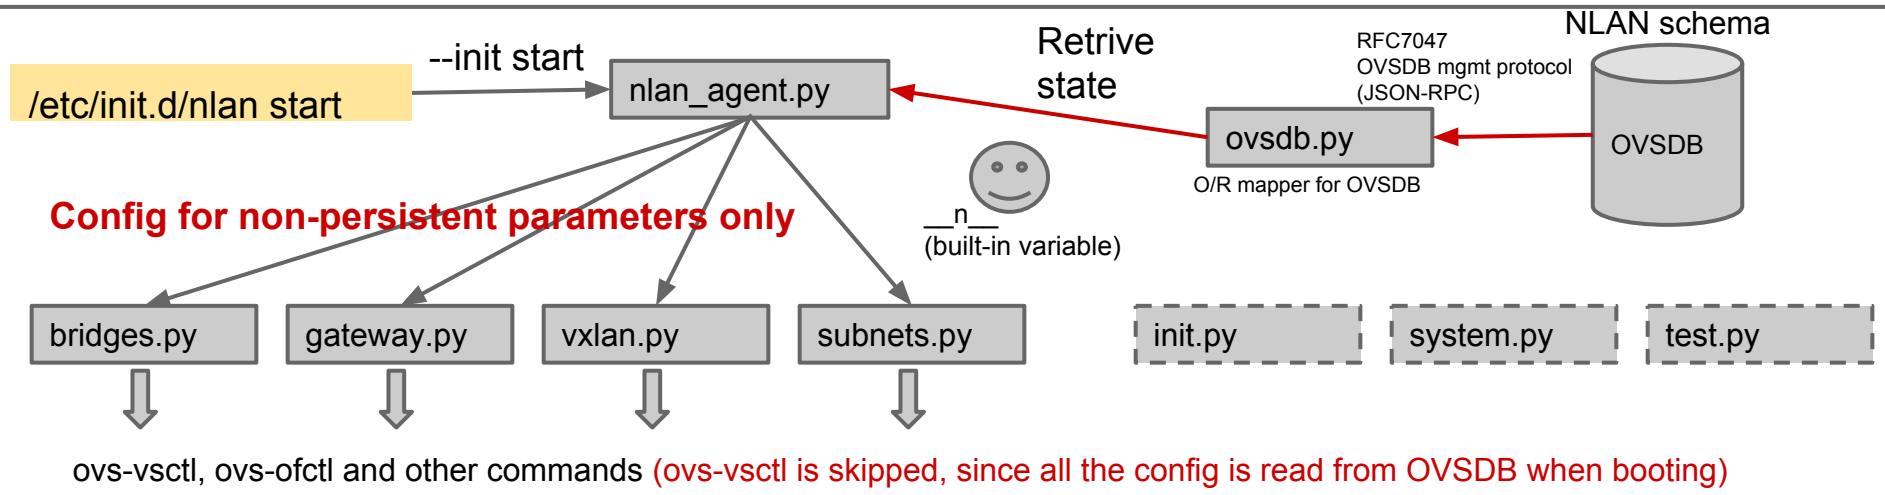
NLAN architecture

(System reboot)

DevOps master (Debian)

(All the routers are independent from the master)

Router (OpenWrt or Raspbian)



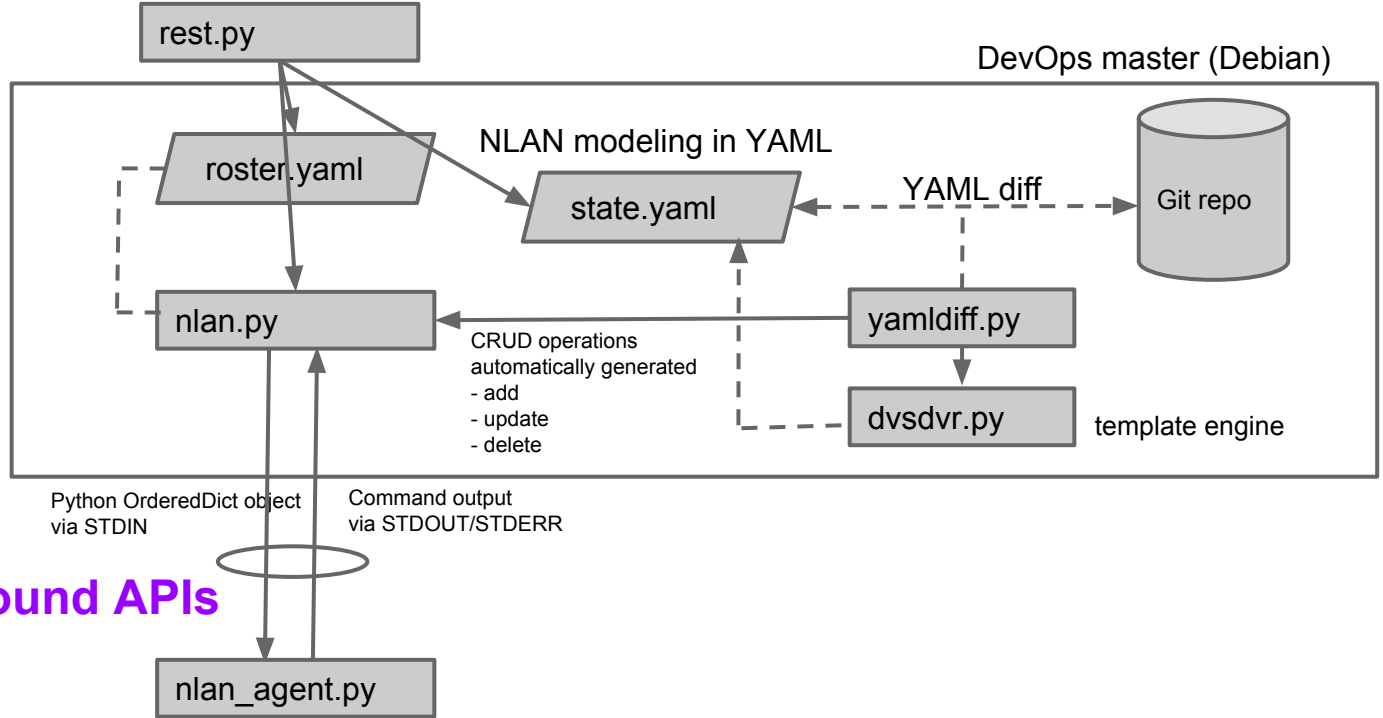
NLAN architecture (REST APIs)

Northbound APIs

It's a bit like RESTCONF...



HTTP GET/POST/PUT/DELETE/OPTIONS w/ query parameters



OpFlex -- Flexible OVSDB schemas

OpFlex:

<http://www.cisco.com/c/en/us/solutions/collateral/data-center-virtualization/application-centric-infrastructure/white-paper-c11-731302.html>

OpFlex:

“It uses dynamic, flexible schemas for interaction with devices, effectively increasing the network to a higher common denominator feature set. The Open vSwitch Database (OVSDB) management protocol allows configuration of high-level abstract data models as well as basic primitives such as ports and bridges, and can support SDN geeks’ innovations”

OVSDB schema (Open_vSwitch database)

NLAN schema in YAML

```
      :
NLAN_Subnet:
  columns:
  vni:
    type:
    key: {type: integer, minInteger: 0, maxInteger: 16777215}
    min: 1
    max: 1
    _description: "Virtual network identifier"
  vid:
    type:
    key: {type: integer, minInteger: 0, maxInteger: 4095}
    min: 0
    max: 1
    _description: "VLAN ID"
  ip_dvr:
    type:
    key: {type: string, enum: [set, [addr, mode, dhcp]]}
    value: {type: string, _pattern: {addr: ipv4_prefix, mode: dvr_mode, dhcp:
string}}
    min: 0
    max: 3
    _description: "Distributed Virtual Router setting"
  ip_vhost:
    type:
    key: {type: string, _pattern: ipv4_prefix}
    min: 0
    max: 1
    _description: "Virtual host in a linux network namespace"
  default_gw:
    type:
    key: {type: string, _pattern: ipv4_address}
    min: 0
    max: 1
    _description: "Default GW address for this subnet"
      :
```

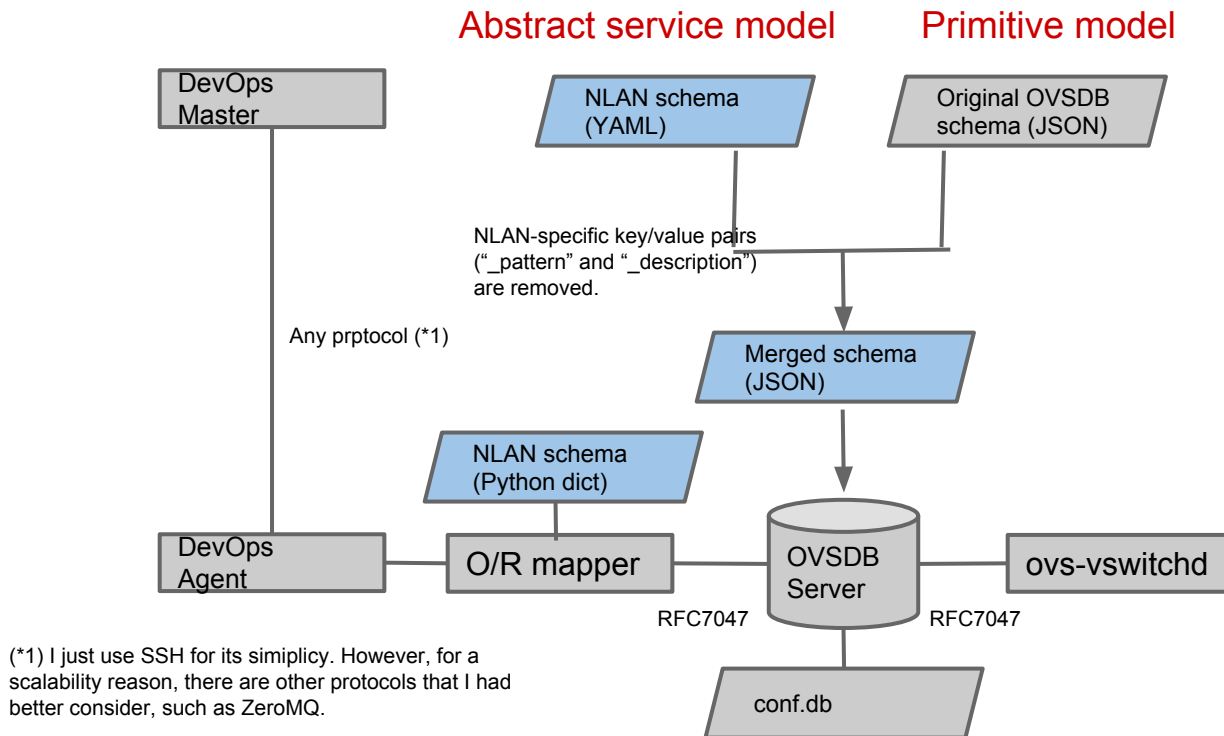
Original OVSDB schema in JSON

```
{
  "name": "Open_vSwitch",
  "version": "7.4.2",
  "cksum": "951746691 20389",
  "tables": {
    "NLAN": {
      "columns": {
        "bridges": {
          "type": {
            "key": {
              "type": "uuid",
              "refTable": "NLAN_Bridges",
              "min": 0,
              "max": 1
            },
            "services": {
              "type": {
                "key": {
                  "type": "uuid",
                  "refTable": "NLAN_Service",
                  "min": 0,
                  "max": "unlimited"
                },
                "gateway": {
                  "type": {
                    "key": {
                      "type": "uuid",
                      "refTable": "NLAN_Gateway",
                      "min": 0,
                      "max": 1
                    },
                    "vxlan": {
                      "type": {
                        "key": {
                          "type": "uuid",
                          "refTable": "NLAN_VXLAN",
                          "min": 0,
                          "max": 1
                        },
                        "subnets": {
                          "type": {
                            "key": {
                              "type": "uuid",
                              "refTable": "NLAN_Subnet",
                              "min": 0,
                              "max": "unlimited"
                            },
                            "isRoot": true,
                            "maxRows": 1
                          },

```

Merging schemas

OVSDB schema can also express more abstract data models.



NLAN states in OVSDDB

(ovsdb-client dump Open_vSwitch)

```
NLAN_table
_uuid          bridges          gateway          services subnets
              vxlan
-----
f41f2732-c9a0-42ef-a72f-3275ec79c13d 90af6783-f57c-418f-8102-6e02e30d1427 1b3f608c-9a3e-4c4a-9b7c-9805cb8d0d69 [] [3965d784-b309-4d64-b700-396d1ad0a2a4, 53cc4e84-d7b7-4a1e-9d02-df29fa5d9afc, aa54b72e-a22c-47a6-8c0f-c537da6fa8cf] 88d7ad5f-158f-437e-9169-78f79d2e38d5
```

```
NLAN_Bridges_table
_uuid          controller ovs_bridges
-----
90af6783-f57c-418f-8102-6e02e30d1427 [] enabled
```


```
NLAN_Gateway_table
_uuid          network rip
-----
1b3f608c-9a3e-4c4a-9b7c-9805cb8d0d69 "eth2" enabled
```

```
NLAN_Service_table
_uuid chain name
-----
```

```
NLAN_Subnet_table
_uuid          default_gw ip_dvr          ip_vhost          peers          ports vid vni
-----
53cc4e84-d7b7-4a1e-9d02-df29fa5d9afc [] {addr="10.0.1.1/24", mode=hub} "10.0.1.101/24" ["192.168.1.104"] [] 1 1001
3965d784-b309-4d64-b700-396d1ad0a2a4 [] {addr="10.0.3.1/24", mode=dvr} "10.0.3.101/24" ["192.168.1.102", "192.168.1.103"] [] 3 103
aa54b72e-a22c-47a6-8c0f-c537da6fa8cf [] {addr="192.168.100.1/24", mode=dvr} "192.168.100.101/24" ["192.168.1.102", "192.168.1.103"] [] 2 1
```

```
NLAN_VXLAN_table
_uuid          local_ip          remote_ips
-----
88d7ad5f-158f-437e-9169-78f79d2e38d5 "192.168.1.101" ["192.168.1.102", "192.168.1.103", "192.168.1.104"]
```

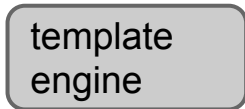
Model-driven service abstraction (cont'd)




Step1: define network service model



Step2: write the mode as "desired state" in YAML format w/ some placeholders for a template engine

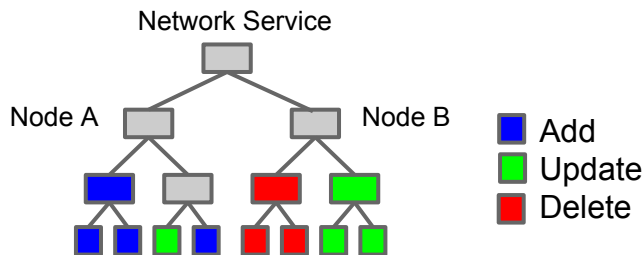


Step3: write a template engine to fill out the placeholders automatically.



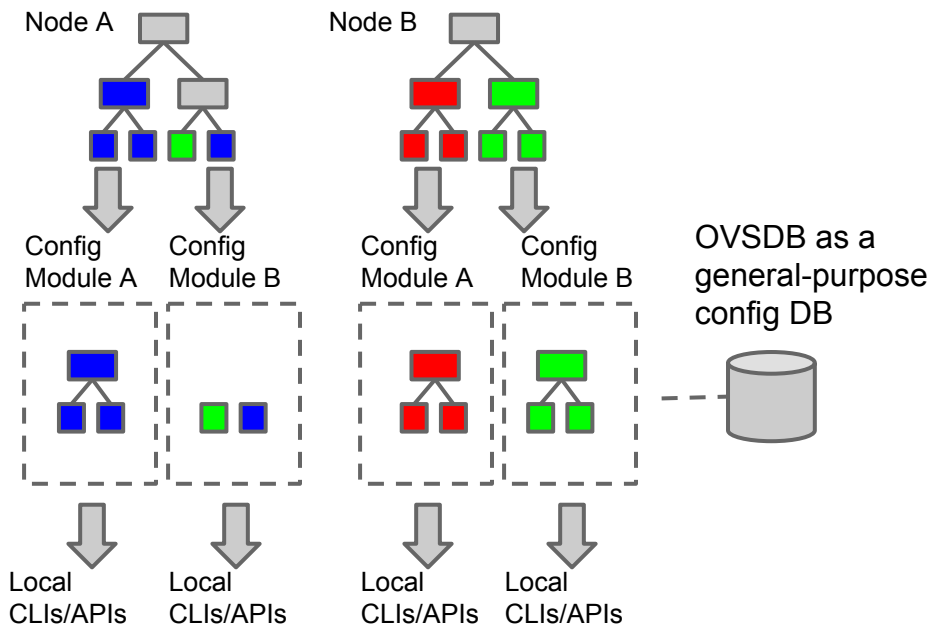
Step4: NLAN-Master (nlan.py and yamldiff.py) generates CRUD operations comparing the desired state with the previous state

Model-driven service abstraction



Step5: Now CRUD-operations (= diff between the previous and the desired states) are in the form of Python OrderedDict object

Step6: NLAN-Agent (nlan_agent.py) routes the CRUD operations to corresponding nodes/modules



Template and placeholders example

`#!/template.dvsvdvr`

`openwrt1:`

`:`

`vxlan:`

`local_ip: <local_ip>`

`remote_ips: <remote_ips>`

`subnets:`

`- vid: 1`

`vni: 101`

`ip_dvr: {addr: '10.0.1.1/24', mode: dvr}`

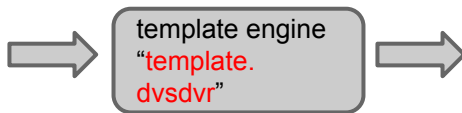
`ip_vhost: '10.0.1.101/24'`

`ports:`

`- eth0.1`

`peers: <peers>`

`:`



- generates a local ip address
- generates VXLAN remote ip addresses
- generates broadcast tree per VNI
- automatically resolves dependencies among parameters

`openwrt1:`

`:`

`vxlan:`

`local_ip: '192.168.1.101'`

`remote_ips: ['192.168.1.102', '192.168.1.103', '192.168.1.104']`

`subnets:`

`- vid: 1`

`vni: 101`

`ip_dvr: {addr: '10.0.1.1/24', mode: dvr}`

`ip_vhost: '10.0.1.101/24'`

`ports:`

`- eth0.1`

`peers: ['192.168.1.102', '192.168.1.104']`

`:`

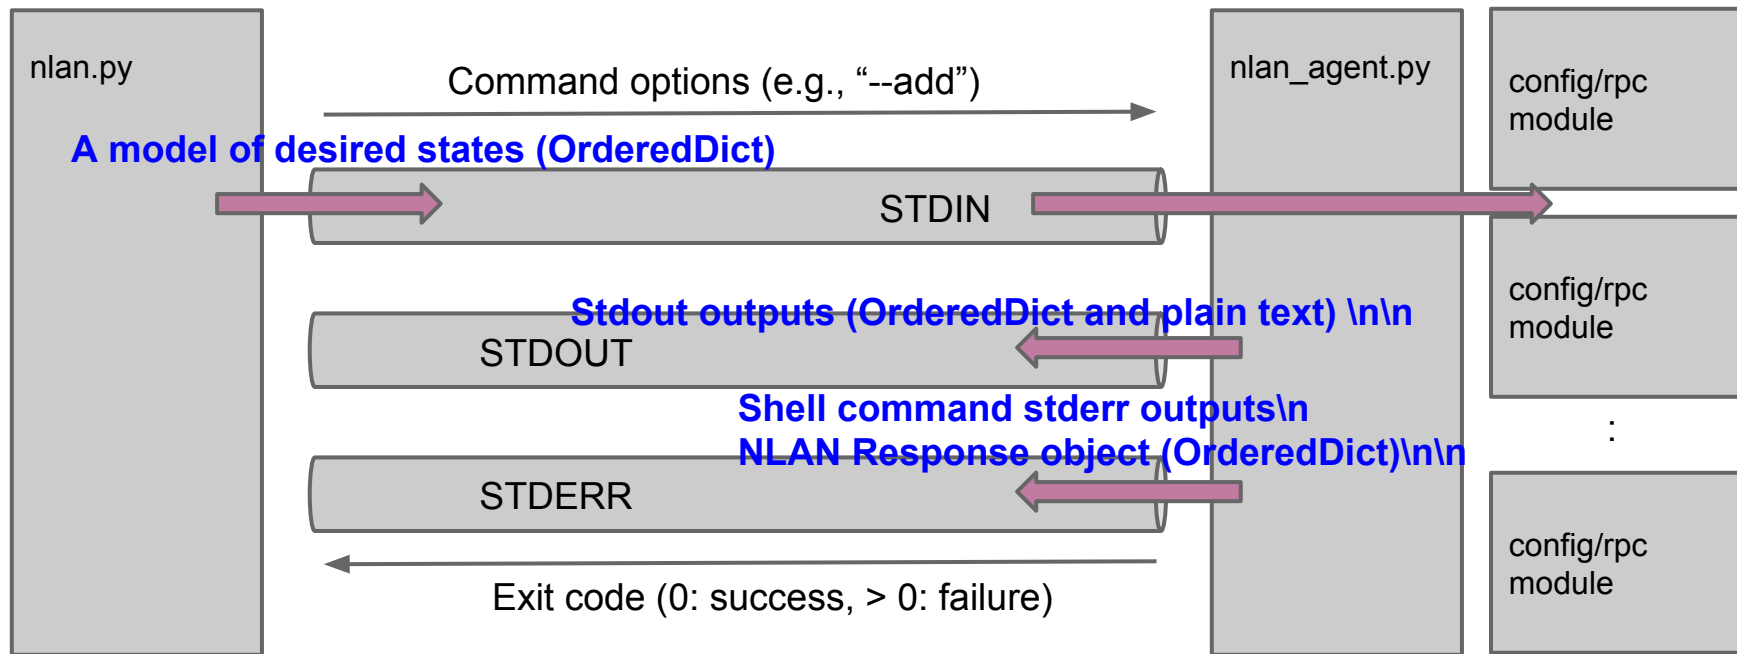
Model of desired state

- NLAN-Master sends Python OrderedDict to NLAN-Agent via ssh STDIN.
- To be exact, string form of an OrderedDict object (sort of object serialization).

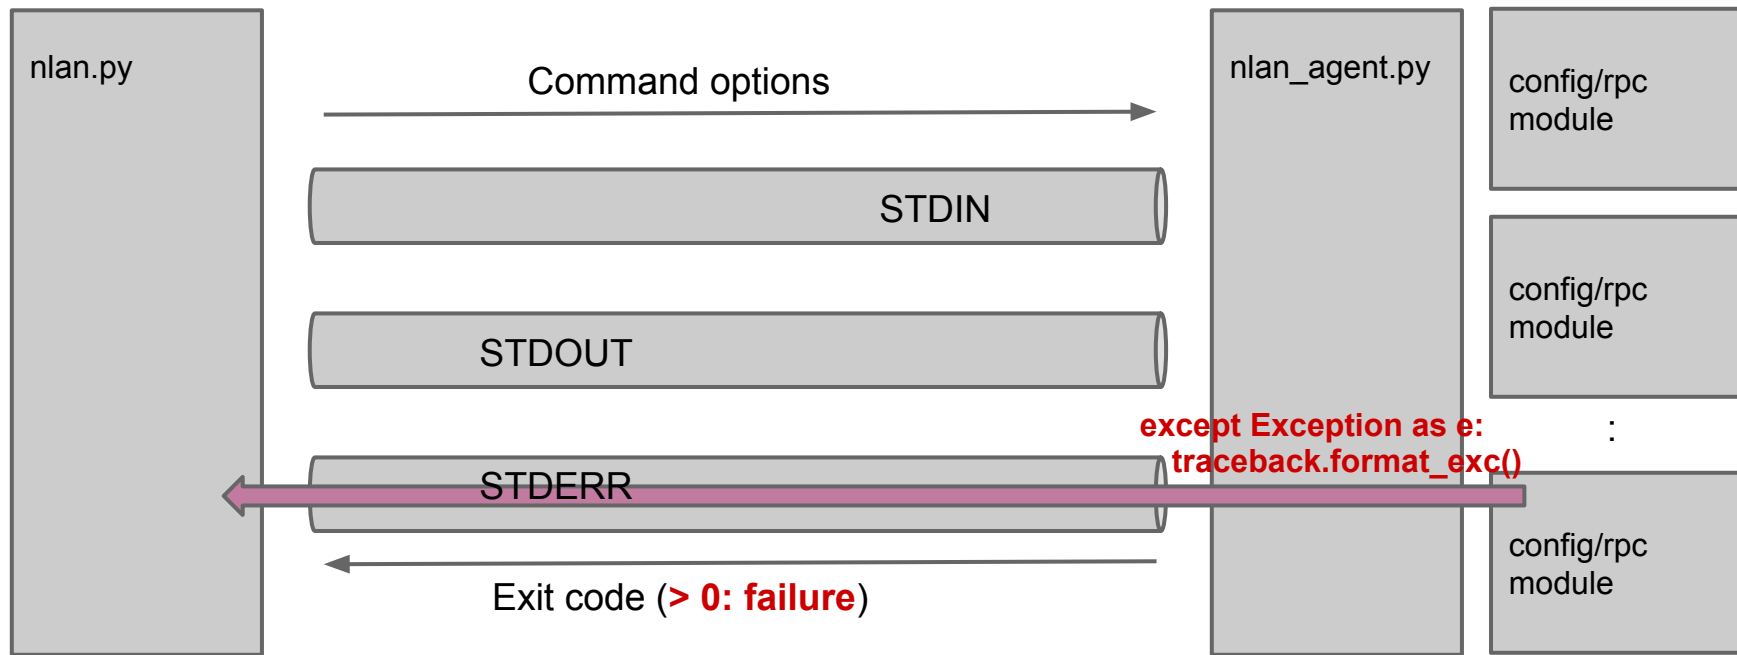
```
"OrderedDict([('bridges', {'ovs_bridges': 'enabled'}), ('gateway', {'network': 'eth2', 'rip':  
'enabled'}), ('vxlan', {'remote_ips': ['192.168.1.103', '192.168.1.102', '192.168.1.104'], 'local_ip':  
'192.168.1.101'}), ('subnets', [{'peers': ['192.168.1.102', '192.168.1.103'], 'vid': 2, '_index': ['vni', 1],  
'ip_vhost': '192.168.100.101/24', 'vni': 1, 'ip_dvr': OrderedDict([('addr', '192.168.100.1/24'),  
('mode', 'dvr')])}], {'peers': ['192.168.1.102', '192.168.1.103'], 'vid': 3, '_index': ['vni', 103],  
'ip_vhost': '10.0.3.101/24', 'vni': 103, 'ip_dvr': OrderedDict([('addr', '10.0.3.1/24'), ('mode', 'dvr')])}],  
{'peers': ['192.168.1.104'], 'vid': 1, '_index': ['vni', 1001], 'ip_vhost': '10.0.1.101/24', 'vni': 1001,  
'ip_dvr': OrderedDict([('addr', '10.0.1.1/24'), ('mode', 'hub')])}]])"
```

- Imperative/declarative state representation.
- I don't use JSON, since NLAN is 100% Python implementation.

NLAN Request/Response over SSH

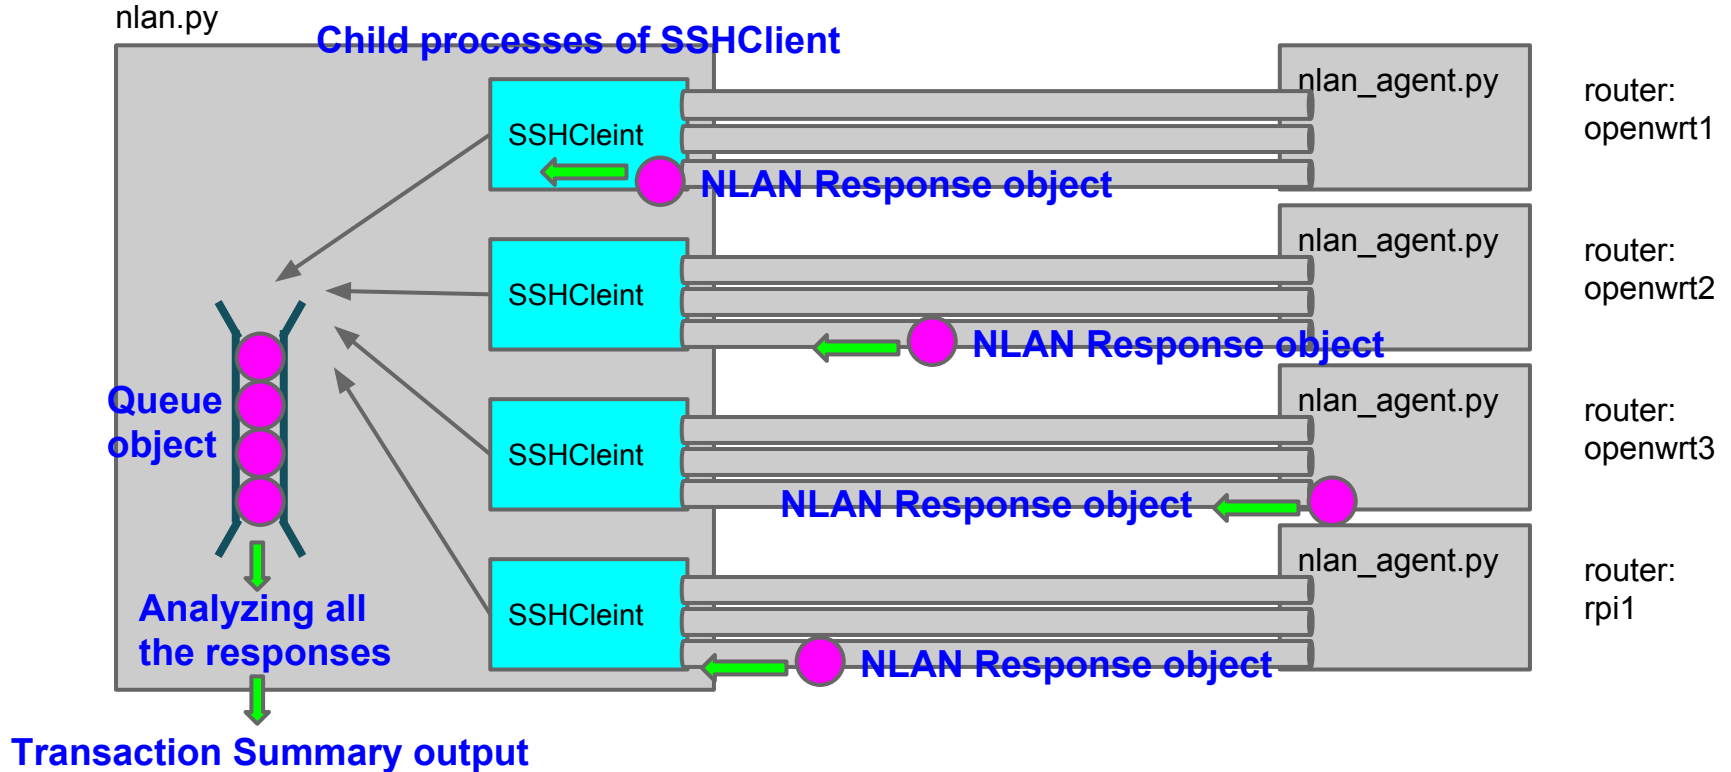


Exception handling



Trouble shooting becomes much easier by sending `traceback.format_exc()` and raw command outputs (plain text via stdout/stderr) to `nlan.py`.

Parallel SSH sessions



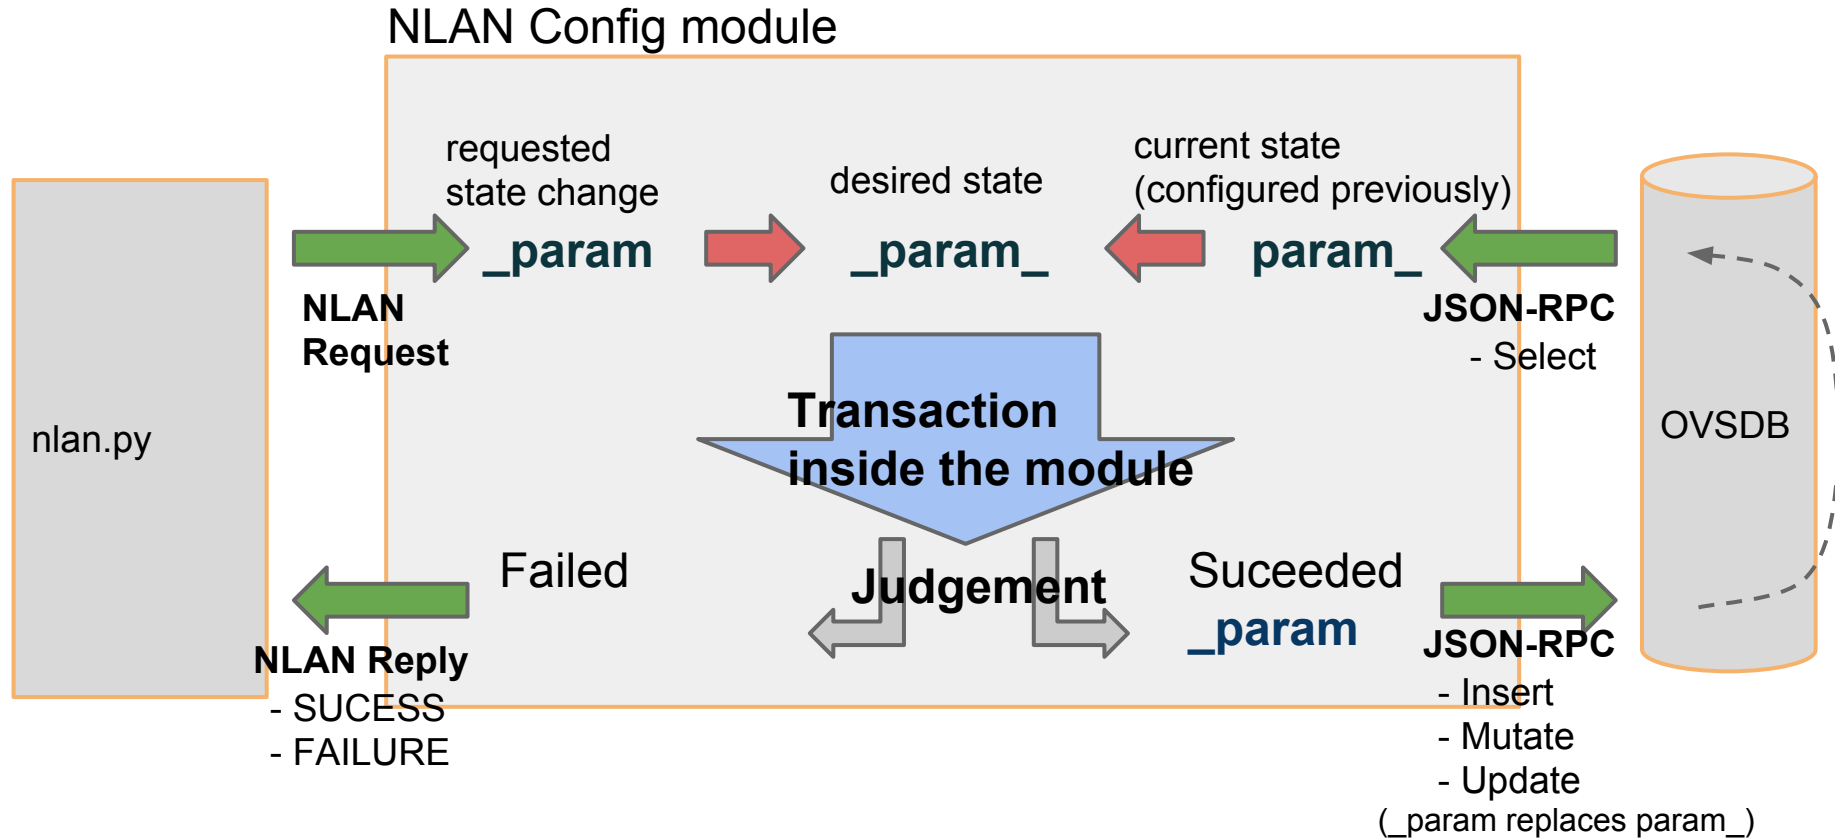
Transaction Summary output

Transaction Summary

```
Start Time: 2014-04-22 20:48:31.835552
```

Router	Result	Elapsed Time
openwrt1	: -)	2.88 (sec)
openwrt3	: -)	2.99 (sec)
openwrt2	: -)	3.00 (sec)
rpil	: -)	3.08 (sec)

CRUD operations inside NLAN config modules



Global variables (__dict__) generated by CRUD.params()

model

```
state +- param a
|
+- param b
|
+- param c
|
+- param d
```

State parameters (for add/update operations)

			Global variables generated by Model.params()			
State params	Requested by Master	OVSDB	_param (Requested change)	_param_ (Desired state)	param_ (Current state in OVSDB)	Operation
a	1	None	_a=1	_a_=1	a_=None	add
b	2	1	_b=2	_b_=2	b_=1	update
c	None	1	_c=None	_c_=1	c_=1	
d	None	None	_d=None	_d_=None	d_=None	

State parameters (for delete operations)

			Global variables generated by Model.params()			
State params	Requested by Master	OVSDB	_param (Requested change)	_param_ (Desired state)	param_ (Current state in OVSDB)	Operation
a	1	None	_a=1	_a_=***	a_=None	(Never exists)
b	1	1	_b=1	_b_=None	b_=1	delete
c	None	1	_c=None	_c_=1	c_=1	
d	None	None	_d=None	_d_=None	d_=None	

Python coding in NLAN Config modules

nlan_agent.py

```
with oputil.CRUD(...):  
    module.add()
```



module.py

```
def add():  
    if _param1:  
        (execute local commands w/ _param1 and other  
         _param_(s) as arguments)  
    if _param2:  
        (execute local commands w/ _param2 and other  
         _param_(s) as arguments)  
    if _param3:  
        (execute local commands w/ _param3 and other  
         _param_(s) as arguments)
```

```
with oputil.CRUD(...):  
    module.delete()
```



```
def delete():
```

: Every NLAN config
module MUST implement
add(), delete() and
update() functions.


```
with oputil.CRUD(...):  
    module.update()
```



```
def update():
```

Service Function Chaining in YAML

```
rpi1:
  bridges:
    ovs_bridges: enabled
  services: # Service Functions
    - name: snort1
      chain: [mz.101, dmz.1001]
  vxlan:
    local_ip: <local_ip>
    remote_ips: <remote_ips>
  subnets:
    - vid: 111
      vni: 1001
      peers: <peers>
      ports: <sfports>
    - vid: 1
      vni: 101
      peers: <peers>
      ports: <sfports>
```



NLAN command usage (nlan.py)

Copy NLAN-agent-side modules to all the target routers (incl. NLAN/OVSDB schema and Linux init.d scripts):

```
$ nlan.py --scpmmod
```

Initialize states at all the target routers:

```
$ nlan.py init.run
```

Ask all the target routers to transit to the desired states

```
$ nlan.py -G deploy
```

Rollback to the previous config

```
$ nlan.py init.run
```

```
$ nlan.py -R deploy
```

Command line CRUD (add/get/update/delete) operations

```
$ nlan.py -t openwrt1 --add subnets _index=101 vid=1 ip_dvr=mode:dvr,addr:10.0.1.1/24 ip_vhost=10.0.1.101/24
```

```
$ nlan.py -t openwrt1 --update subnets _index=101 ip_vhost=10.0.1.109/24
```

```
$ nlan.py -t openwrt1 --delete subnets _index=101 ip_vhost=10.0.1.109/24
```

```
$ nlan.py -t openwrt1 --get subnets _index=101 vid ip_dvr
```

Reboot all the target routers:

```
$ nlan.py system.reboot
```

NLAN command usage (nlan_agent.py)

Initialize states:

```
$ nlan_agent.py init.run
```

Command line CRUD (add/get/update/delete) operations

```
$ nlan_agent.py --add subnets _index=101 vid=1 ip_dvr=mode:dvr,addr:10.0.1.1/24 ip_vhost=10.0.1.101/24
```

```
$ nlan_agent.py --update subnets _index=101 ip_vhost=10.0.1.109/24
```

```
$ nlan_agent.py --delete subnets _index=101 ip_vhost=10.0.1.109/24
```

```
$ nlan_agent.py --get subnets _index=101 vid ip_dvr
```

Reboot all the target routers:

```
$ nlan.py system.reboot
```

REST APIs


HTTP Method	NLAN CRUD operations	NLAN RPC operations
GET	CRUD: get URL: /<router>/ config /<module>/<_index>	URL: /<router>/ rpc /<module>/<func>
POST	CRUD: add URL: /<router>/ config /<module>/<_index>	-
PUT	CRUD: update URL: /<router>/ config /<module>/<_index>	-
DELETE	CRUD: delete URL: /<router>/ config /<module>/<_index>	-
OPTIONS	Get NLAN schemas URL: none	-

REST APIs example

HTTP Method	URL	Query parameters
POST	/_ALL/rpc/test/echo	params=Hello!
OPTIONS	(none)	params=subnets
POST	/openwrt1/rpc/init/run	(none)
POST	/openwrt1/config/bridges	ovs_bridges=enabled
POST	/openwrt1/config/vxlan	local_ip=192.168.1.101&remote_ips=192.168.1.102,192.168.56.103
PUT	/openwrt1/config/vxlan	remote_ips=192.168.1.102,192.168.56.104
GET	/openwrt1/config/vxlan	params=remote_ips
POST	/openwrt1/config/subnets/101	vni=101&vid=1&ip_dvr=addr:10.0.1.1/24,mode:dvr
DELETE	/openwrt1/config/subnets/101	params=ip_dvr
POST	/openwrt1/rpc/db/state	(none)

Scenario Runner

- nlans.py -- reads test scenarios and executes each test
- Test scenarios written in YAML
- Automatic test result confirmation
 - Inspired by Python's "unittest"
 - "assert"
 - "asserRaises"
 - ("assertOutputs" to be supported)



Simple RPC library

```
import rpc
```

```
# calls test.kwargs_test(...) at router 'openwrt1'
```

```
rpc = rpc.RPC(module='test', func='kwargs_test', target='openwrt1')
```

```
result = rpc(1, b='Hello', d='World!')
```

```
print result
```

```
# calls test.echo(*args) at all routers on the roster
```

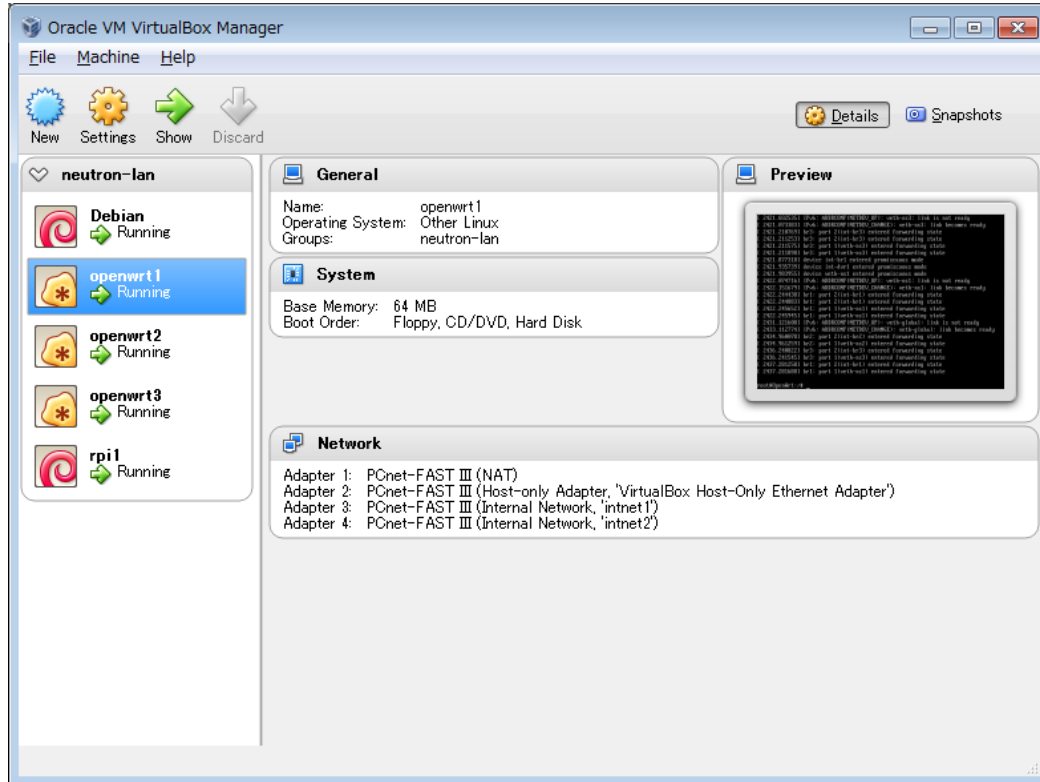
```
rpc = rpc.RPC(module='test', func='echo')
```

```
results = rpc('Hello World!')
```

```
for l in results:
```

```
    print l['router'], l['stdout']
```

NLAN Software Development environment on VirtualBox

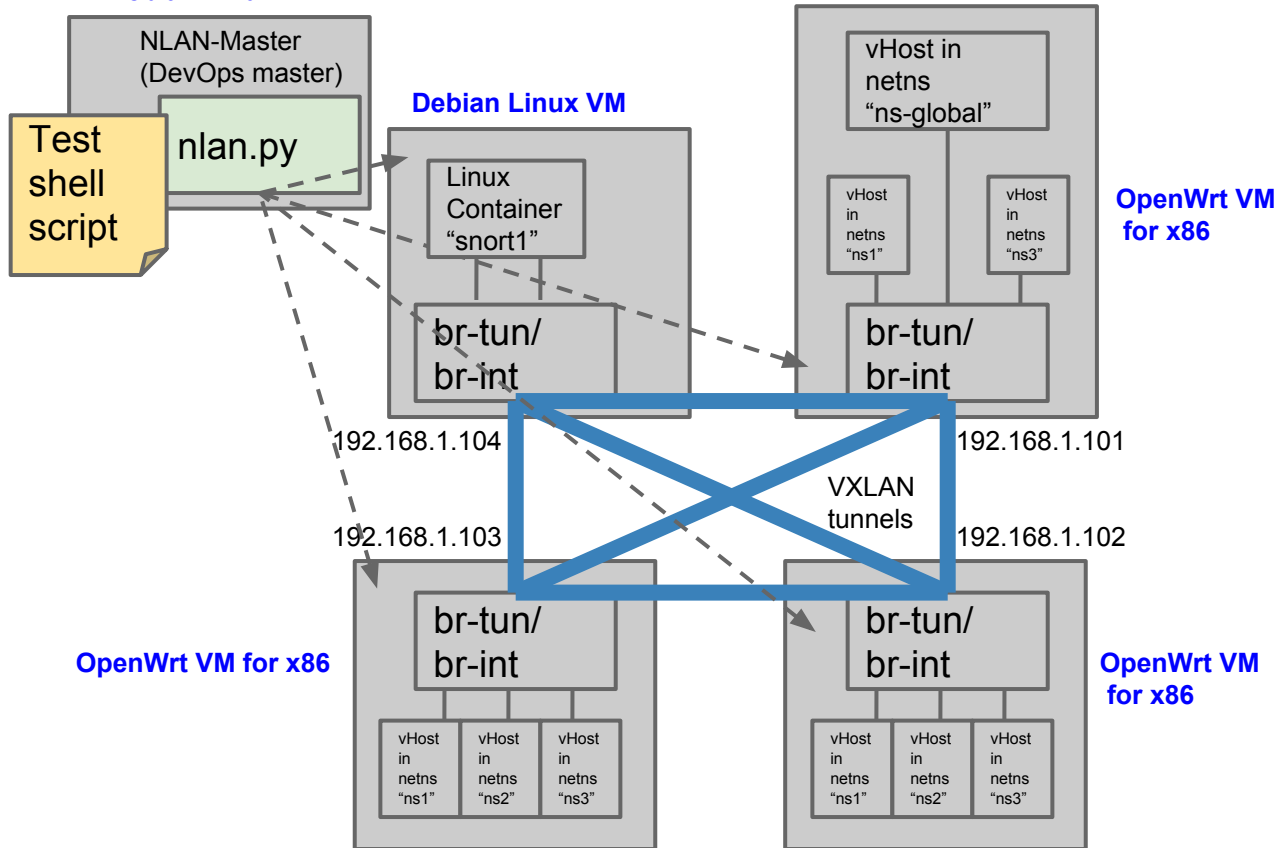


- Five VMs running on one Win7 PC.
 - Two Debian VMs
 - Three OpenWrt VMs
- OpenWrt image for x86
 - I built the kernel with Open vSwitch 2.0.0 and netns/veth/LXC support
 - Very light-weight Linux supporting Open vSwitch 2.0.0 ⇒ An alternative to mininet 2.0.0
- Network adapters setting
 - Internet access: "NAT"
 - Management: "Host-Only"
 - NLAN underlay: "Internal"

Integration Test environment on VirtualBox

(running the test script every day)

Debian Linux VM



- Open vSwitch-based network more realistic than mininet 2.0
 - Every vSwitch with full-fledged(?) Linux
 - netns-based virtual hosts
- Mimics "Beremetal Switch"
- Integration Test script running on Debian Linux VM
 - Makes use of "nlan.py"